

Listen

Learning Doc 03: Commands and the Intent Engine
How “search what is an atom” becomes an action

Detailed design, Section 2

7 October 2026


Abstract

The voice module hands us a line of text, such as "**search what is an atom**". This document designs everything between that text and an action on your screen: the file format that describes commands, the matching algorithm that picks the right command, how default apps work, how you add and edit commands without code, and why the design is safe even when an AI writes the commands. Each new idea is introduced as a concept box.

Contents

1	The problem in one picture	2
2	The command file	2
2.1	Data-driven design	2
2.2	Three sections: targets, defaults, commands	2
2.3	Slot types	3
3	Step 1: normalisation	4
4	Step 2: matching	4
4.1	Patterns become regular expressions	4
4.2	When several patterns match: specificity	4
4.3	Validation can reject a match	5
4.4	Defaults fill missing slots	5
4.5	No match is useful information	5
4.6	The whole algorithm	5
5	Finding apps by name	5
6	Step 3: the contract with the action executor	6
6.1	The v1 action primitives	6
7	Adding and editing commands without code	6
7.1	Where the file lives	6
7.2	Validation on every load	7
7.3	The command builder	7
7.4	Why this is safe	7
8	Testing this design	7
9	Glossary	8
10	Check your understanding	8

1 The problem in one picture



Three steps: **clean** the text, **match** it to a command and pull out the interesting parts, then **act**. This document covers steps 1 and 2 in detail, and the contract between step 2 and step 3.

Concept: Intent and slots

An **intent** is *what you want done* (“search”). **Slots** are the *details* needed to do it (the query “what is an atom”, the target “google”). Turning a sentence into an intent plus filled slots is called **intent parsing** or **slot filling**. Siri, Alexa and Google Assistant all have this exact step.

2 The command file

2.1 Data-driven design

Concept: Data-driven design

Instead of writing behaviour as code (if `text.startswith("search")...`), you describe behaviour in a **data file**, and write one small, general piece of code that reads the file and follows it. Changing behaviour means editing data, not code.

This is what makes your requirement “add play {query} on spotify without touching code” possible. The file is written in **YAML**, a format designed to be easy for humans to read and edit: indentation shows nesting, `-` starts a list item, `key: value` sets a field.

2.2 Three sections: targets, defaults, commands

```

version: 1

# targets: the places Listen can send you
targets:
  google: { search: "https://www.google.com/search?q={query}",
             home:  "https://www.google.com" }
  youtube: { play:  "https://www.youtube.com/results?search_query={query}",
             search: "https://www.youtube.com/results?search_query={query}",
             home:  "https://www.youtube.com" }
  spotify: { play:  "https://open.spotify.com/search/{query}",
             home:  "https://open.spotify.com" }

# defaults: which target a verb uses when you don't name one
defaults:
  search: google
  play:  youtube

# commands: sentence patterns -> an action primitive
commands:
  - id: search
    patterns: ["search {query} on {target}", "search {query}",

```

```

        "look up {query}", "google {query}"]
    action: open_url
    url_from: search          # use the target's "search" URL

- id: play
  patterns: ["play {query} on {target}", "play {query}"]
  action: open_url
  url_from: play

- id: open
  patterns: ["open {app}", "launch {app}", "start {app}"]
  action: launch_app

- id: close
  patterns: ["close {app}", "quit {app}", "exit {app}"]
  action: close_app

- id: volume_set
  patterns: ["volume {number}", "set volume to {number}"]
  action: set_volume

- id: volume_up
  patterns: ["volume up", "louder", "turn it up"]
  action: change_volume
  params: { step: 10 }

- id: volume_down
  patterns: ["volume down", "quieter", "turn it down"]
  action: change_volume
  params: { step: -10 }

- id: mute
  patterns: ["mute", "unmute"]
  action: toggle_mute

```

Why it is done this way

Targets are separated from **commands** because “where” and “what” change for different reasons (the single-responsibility idea from Doc 02). Look at what this buys you: `play {query} on spotify` needs *no new command at all*. The `play` command already accepts `on {target}`; you only add a `spotify` target with a `play` URL. One new line of data, and *every* verb that has a URL for Spotify starts working with it.

2.3 Slot types

Every `{name}` in a pattern is a slot, and its name decides what it may contain:

Slot	Accepts	Example
<code>{query}</code>	Any text (at least one word)	“what is an atom”
<code>{target}</code>	Only names in targets	“spotify”
<code>{app}</code>	An app Listen can find on your PC (Section 5)	“vs code”
<code>{number}</code>	A number, written as digits or words	“50”, “fifty”

Typed slots are how the matcher tells “play *songs on the beach*” (the *query* is “songs on the beach”, because “the beach” is not a target) apart from “play *lofi on spotify*”.

3 Step 1: normalisation

Concept: Normalisation

Different-looking inputs that mean the same thing are converted into one standard form *before* comparing them. “Search what is an atom.”, “search what is an atom” and “SEARCH what is an atom?” must all match the same pattern.

The cleaner, in order:

1. Lowercase everything.
2. Remove punctuation the speech engine adds (. , ? !), but keep characters that matter inside queries such as + and ' in “what’s”.
3. Remove the wake word if it leaked into the transcript (“hey jarvis search. .”).
4. Remove polite filler at the start and end: “please”, “can you”, “could you”, “for me”. (“Can you please search atoms for me” → “search atoms”).
5. Collapse repeated spaces.

4 Step 2: matching

4.1 Patterns become regular expressions

Concept: Regular expression (regex)

A **regex** is a tiny language for describing text patterns. For example `^search (?P<query>.+)$` means: the text *starts* (^) with “search ”, then one or more characters (.⁺) that we capture under the name `query`, then the text *ends* (\$). Python’s built-in `re` module runs them.

You never write regexes yourself. When the command file loads, each friendly pattern is *compiled* into one:

Pattern in YAML	Regex the engine builds
<code>search {query}</code>	<code>^search (?P<query>.+)?\$</code>
<code>play {query} on {target}</code>	<code>^play (?P<query>.+?) on (?P<target>.+)?\$</code>

The `^` and `$` *anchors* mean the whole sentence must fit the pattern, so “research papers” does not accidentally match “search {query}”.

4.2 When several patterns match: specificity

“play lofi on spotify” matches *both*

- `play {query} on {target}` with query = “lofi”, target = “spotify”, and
- `play {query}` with query = “lofi on spotify”.

Concept: Specificity

When several rules match, prefer the **most specific** one: the one that explains more of the sentence with fixed words. Score = number of literal words in the pattern. The first pattern has 2 (“play”, “on”); the second has 1. The first wins. On a tie, the command listed earlier in the file wins, so the order of the file is a tiebreaker you control.

4.3 Validation can reject a match

After a regex matches, each slot is checked against its type. If a check fails, *that candidate is thrown away* and the next-best one is tried:

```
"play songs on the beach"
  candidate 1: play {query} on {target}  target="the beach" -> not a target, REJECT
  candidate 2: play {query}              query="songs on the beach" -> ACCEPT
```

4.4 Defaults fill missing slots

If the winning pattern has no `{target}` (“play lofi”), the engine looks up `defaults.play` → `youtube`. If the target has no URL for that verb (asking Google to “play” something), the result is a clear error message instead of a guess.

4.5 No match is useful information

If nothing matches, the engine publishes an `IntentNotMatched` event. The terminal says “Sorry, I didn’t understand: ...”, and the sentence is logged. Later, your analytics page can list *what you tried that failed*: exactly the commands worth adding next.

4.6 The whole algorithm

```
match(text):
    text = normalise(text)
    candidates = []
    for each command, for each pattern:
        if pattern.regex fully matches text:
            slots = the captured groups
            if every slot passes its type check:
                candidates.append( (specificity, file_order, command, slots) )
    if no candidates: return NoMatch(text)
    best = highest specificity, then earliest in file
    fill missing target from defaults
    return Intent(command, slots)
```

Why it is done this way

This is a **pure function**: text in, result out, no microphone, no files, no internet, no side effects. That makes it trivial to test with a table of example sentences, and it runs in well under a millisecond, so matching costs nothing on your lightweight budget.

5 Finding apps by name

“open vs code” requires knowing that “vs code” means `C:\...\Code.exe`. Windows already keeps that knowledge: every installed app puts a **shortcut** in the Start Menu folders.

1. At startup the agent scans the Start Menu shortcut folders and builds an **app index**: display name → how to launch it. This takes a fraction of a second and is refreshed when you ask (“refresh apps”) or once a day.
2. You can add **aliases** in the command file (`"code": "Visual Studio Code"`) for names the speech engine spells oddly.
3. The `{app}` slot is matched with **fuzzy matching**.

Concept: Fuzzy matching

Exact comparison fails on “vs code” vs. “Visual Studio Code” or “spotify” vs. “Spotify”. Fuzzy matching gives a **similarity score** between 0 and 1 and accepts the best candidate above a threshold (we start at 0.8 and tune it). Python’s built-in `difflib` does this with no extra download.

“open google” works too: if the name is not an app but *is* a target, Listen opens that target’s home URL.

6 Step 3: the contract with the action executor

The intent engine never performs actions itself. It publishes an `IntentMatched` event:

```
IntentMatched(
    command_id = "search",
    action      = "open_url",
    args       = { "url": "https://www.google.com/search?q=what+is+an+atom" },
    heard      = "search what is an atom",
)
```

The `actions/` module subscribes and runs it. Values placed into URLs are **URL-encoded** (spaces become +, & becomes %26), so a query can never break the URL or sneak in extra parameters.

6.1 The v1 action primitives

Primitive	What it does
<code>open_url</code>	Opens a web address in your default browser
<code>launch_app</code>	Starts an app from the app index
<code>close_app</code>	Asks an app’s windows to close (like clicking X, so unsaved-work prompts still appear)
<code>set_volume</code>	Sets master volume to 0–100
<code>change_volume</code>	Raises or lowers volume by a step
<code>toggle_mute</code>	Mutes or unmutes

Watch out

Honest limitation: `play lofi` opens YouTube’s *search results* for “lofi”; it does not start the first video automatically. Autoplay would need either the YouTube API or scraping the results page. It is a possible later improvement, deliberately out of v1.

7 Adding and editing commands without code

7.1 Where the file lives

- A **default** command file ships inside the agent’s code (versioned in git).
- On first run it is copied to your user-data folder (`%APPDATA%\Listen\commands.yaml`). That copy is yours: edits go there, and code updates never overwrite it.
- Once sync exists (Section 5 of the design), the cloud holds the master copy, so edits made in the web app reach your laptop.


7.2 Validation on every load

Concept: Schema validation

A **schema** is a machine-checkable description of what valid data looks like: “every command must have an id, at least one pattern, and an action that is one of the six primitives”. Before using the file, the agent checks it against `shared/recipe.schema.json` (the contract from Doc 02).

An invalid entry is **skipped with a clear warning**, and every valid entry still loads. One typo never takes down the whole assistant.

7.3 The command builder



1. You describe the change in plain words: add a target, add a command, change a default, add a pattern, rename or delete something.
 2. The LLM receives your request, the current file, and the schema, and returns a *proposed change* (not a rewritten file).
 3. The proposal is validated against the schema and tried on a few example sentences (“play lofi on spotify” should now match).
 4. You see a readable diff and approve or reject it. Nothing is saved before “yes”.
- Offline, a guided form asks the same questions step by step.

7.4 Why this is safe

Concept: Allowlist

An **allowlist** is a fixed list of things that are permitted; everything else is refused. Listen’s six primitives are an allowlist. A command can only ever *combine* them; there is deliberately **no** “run any program with any arguments” primitive.

Why it is done this way

An AI is writing data that your computer will act on, so you must ask: what is the worst a bad (or tricked) suggestion could do? With the allowlist, the worst case is “opens a wrong website” or “closes an app politely”. Without it, the worst case is “runs any command on your PC”. Combined with your approval step and URL encoding, this keeps the AI helpful but harmless. Good design limits damage *structurally*, rather than hoping nothing goes wrong.

8 Testing this design

The matcher is a pure function, so its tests are a simple table:

"search what is an atom"	-> search, query="what is an atom", target=google
"Can you search atoms please?"	-> search, query="atoms"
"play lofi"	-> play, query="lofi", target=youtube
"play lofi on spotify"	-> play, query="lofi", target=spotify

"play songs on the beach"	-> play, query="songs on the beach", target=youtube
"volume fifty"	-> volume_set, number=50
"research papers"	-> no match
"open vs code"	-> open, app="Visual Studio Code"

Every bug we ever find in matching becomes one more row in this table, so it can never come back unnoticed. That habit is called a **regression test**.

9 Glossary

Allowlist

A fixed list of permitted things; everything else is refused.

Anchors

Regex symbols `^` (start) and `$` (end) forcing a full match.

Data-driven design

Behaviour described in data files, followed by general code.

Fuzzy matching

Comparing text by similarity score instead of exact equality.

Intent / slot

What you want done / the details needed to do it.

Normalisation

Converting equivalent inputs to one standard form before comparing.

Pure function

Output depends only on input; no side effects. Easy to test.

Regex

A mini-language for describing text patterns.

Regression test

A test added for a fixed bug so it can never silently return.

Schema validation

Checking data against a machine-readable description of valid data.

Specificity

Preferring the rule that explains more of the input with fixed words.

URL encoding

Escaping special characters so text can be safely placed inside a URL.

YAML

A human-friendly data format using indentation for structure.

10 Check your understanding

Check yourself

1. You want “search {query} on bing”. What exactly do you add to the file? Do you need a new command?
2. Why does “research papers” not match `search {query}`?
3. Walk “play songs on the beach” through the algorithm. Which candidates are produced, which is rejected, and why?

4. Why are commands in YAML rather than Python `if` statements?
5. The LLM proposes a command whose action is `run_program`. What happens, and which part of the design stops it?
6. Why is the matcher designed as a pure function? Name two benefits.
7. Your web app should show “commands you tried that failed”. Which event provides this?